

Assignment: Different CNN Architectures on MNIST dataset

```
In [1]: from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
from keras.initializers import he_normal
from keras.layers.normalization import BatchNormalization
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Using TensorFlow backend.

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 1s 0us/step

```
In [2]: if K.image_data_format() == 'channels_first':
        x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
        x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
```

```

        input_shape = (1, img_rows, img_cols)
    else:
        x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
        x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
        input_shape = (img_rows, img_cols, 1)

    x_train = x_train.astype('float32')
    x_test = x_test.astype('float32')
    x_train /= 255
    x_test /= 255
    print('x_train shape:', x_train.shape)
    print(x_train.shape[0], 'train samples')
    print(x_test.shape[0], 'test samples')

    # convert class vectors to binary class matrices
    y_train = keras.utils.to_categorical(y_train, num_classes)
    y_test = keras.utils.to_categorical(y_test, num_classes)

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

```

```

In [0]: # this function is used draw Categorical Crossentropy Loss VS No. of epochs plot
def plt_dynamic(x, vy, ty):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    plt.show()

```

CNN with 3 Convolutional layers and kernel size - (3X3)

```

In [4]: # Initialising the model
model3 = Sequential()

```

```
# Adding first conv layer
model3.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=
input_shape))

# Adding second conv layer
model3.add(Conv2D(64, (3, 3), activation='relu'))

# Adding Maxpooling layer
model3.add(MaxPooling2D(pool_size=(2, 2)))

# Adding Dropout
model3.add(Dropout(0.25))

# Adding third conv layer
model3.add(Conv2D(128, (3, 3), activation='relu'))

# Adding Maxpooling layer
model3.add(MaxPooling2D(pool_size=(2, 2)))

# Adding Dropout
model3.add(Dropout(0.25))

# Adding flatten layer
model3.add(Flatten())

# Adding first hidden layer
model3.add(Dense(256, activation='relu', kernel_initializer=he_normal(se
ed=None)))

# Adding Dropout
model3.add(Dropout(0.5))

# Adding output layer
model3.add(Dense(num_classes, activation='softmax'))

# Printing model Summary
print(model3.summary())

# Compiling the model
```

```
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
# Fitting the data to the model
```

```
history = model3.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
conv2d_2 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_1 (Dropout)	(None, 12, 12, 64)	0
conv2d_3 (Conv2D)	(None, 10, 10, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 128)	0
dropout_2 (Dropout)	(None, 5, 5, 128)	0
flatten_1 (Flatten)	(None, 3200)	0
dense_1 (Dense)	(None, 256)	819456

dropout_3 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 10)	2570

```

=====
Total params: 914,698
Trainable params: 914,698
Non-trainable params: 0

```

```

None
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.cast instead.
Train on 60000 samples, validate on 10000 samples
Epoch 1/12
60000/60000 [=====] - 229s 4ms/step - loss: 0.2175 - acc: 0.9307 - val_loss: 0.0392 - val_acc: 0.9871
Epoch 2/12
60000/60000 [=====] - 228s 4ms/step - loss: 0.0667 - acc: 0.9797 - val_loss: 0.0287 - val_acc: 0.9902
Epoch 3/12
60000/60000 [=====] - 227s 4ms/step - loss: 0.0497 - acc: 0.9852 - val_loss: 0.0232 - val_acc: 0.9918
Epoch 4/12
60000/60000 [=====] - 227s 4ms/step - loss: 0.0414 - acc: 0.9873 - val_loss: 0.0177 - val_acc: 0.9944
Epoch 5/12
60000/60000 [=====] - 227s 4ms/step - loss: 0.0358 - acc: 0.9887 - val_loss: 0.0238 - val_acc: 0.9918
Epoch 6/12
60000/60000 [=====] - 227s 4ms/step - loss: 0.0316 - acc: 0.9900 - val_loss: 0.0207 - val_acc: 0.9933
Epoch 7/12
60000/60000 [=====] - 227s 4ms/step - loss: 0.0282 - acc: 0.9911 - val_loss: 0.0255 - val_acc: 0.9907
Epoch 8/12
60000/60000 [=====] - 227s 4ms/step - loss: 0.

```

```
0248 - acc: 0.9921 - val_loss: 0.0184 - val_acc: 0.9937
Epoch 9/12
60000/60000 [=====] - 227s 4ms/step - loss: 0.
0246 - acc: 0.9922 - val_loss: 0.0170 - val_acc: 0.9945
Epoch 10/12
60000/60000 [=====] - 228s 4ms/step - loss: 0.
0213 - acc: 0.9933 - val_loss: 0.0173 - val_acc: 0.9941
Epoch 11/12
60000/60000 [=====] - 228s 4ms/step - loss: 0.
0197 - acc: 0.9940 - val_loss: 0.0170 - val_acc: 0.9947
Epoch 12/12
60000/60000 [=====] - 227s 4ms/step - loss: 0.
0182 - acc: 0.9939 - val_loss: 0.0169 - val_acc: 0.9948
```

```
In [5]: # Evaluating the model
score = model3.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# Plotting Train and Test Loss VS no. of epochs
# list of epoch numbers
x = list(range(1,epochs+1))

# Validation loss
vy = history.history['val_loss']
# Training loss
ty = history.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)

Test score: 0.016924540818311743
Test accuracy: 0.9948
```



CNN with 5 Convolutional layers and kernel size - (5X5)

```
In [6]: # Initialising the model
model5 = Sequential()

# Adding first conv layer
model5.add(Conv2D(10, kernel_size=(5, 5),padding='same',activation='relu',input_shape=input_shape))

# Adding second conv layer
model5.add(Conv2D(20, (5, 5), activation='relu'))

# Adding Maxpooling layer
model5.add(MaxPooling2D(pool_size=(2, 2),padding='same'))

# Adding Dropout
model5.add(Dropout(0.25))

# Adding third conv layer
```

```
model5.add(Conv2D(35, (5, 5),padding='same', activation='relu'))

# Adding Maxpooling layer
model5.add(MaxPooling2D(pool_size=(2, 2),padding='same'))

# Adding Dropout
model5.add(Dropout(0.25))

# Adding fourth conv layer
model5.add(Conv2D(70, (5, 5),padding='same',activation='relu'))

# Adding fifth conv layer
model5.add(Conv2D(80, (5, 5), activation='relu'))

# Adding Maxpooling layer
model5.add(MaxPooling2D(pool_size=(2, 2),padding='same'))

# Adding Dropout
model5.add(Dropout(0.25))

# Adding flatten layer
model5.add(Flatten())

# Adding first hidden layer
model5.add(Dense(256, activation='relu',kernel_initializer=he_normal(se
ed=None)))

# Adding Batch Normalization
model5.add(BatchNormalization())

# Adding Dropout
model5.add(Dropout(0.5))

# Adding output layer
model5.add(Dense(num_classes, activation='softmax'))

# Printing model Summary
print(model5.summary())
```



```
# Compiling the model
model5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Fitting the data to the model
history= model5.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
```

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 28, 28, 10)	260
conv2d_5 (Conv2D)	(None, 24, 24, 20)	5020
max_pooling2d_3 (MaxPooling2D)	(None, 12, 12, 20)	0
dropout_4 (Dropout)	(None, 12, 12, 20)	0
conv2d_6 (Conv2D)	(None, 12, 12, 35)	17535
max_pooling2d_4 (MaxPooling2D)	(None, 6, 6, 35)	0
dropout_5 (Dropout)	(None, 6, 6, 35)	0
conv2d_7 (Conv2D)	(None, 6, 6, 70)	61320
conv2d_8 (Conv2D)	(None, 2, 2, 80)	140080
max_pooling2d_5 (MaxPooling2D)	(None, 1, 1, 80)	0
dropout_6 (Dropout)	(None, 1, 1, 80)	0
flatten_2 (Flatten)	(None, 80)	0
dense_3 (Dense)	(None, 256)	20736
batch_normalization_1 (Batch Normalization)	(None, 256)	1024

dropout_7 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 10)	2570

Total params: 248,545
 Trainable params: 248,033
 Non-trainable params: 512

None
 Train on 60000 samples, validate on 10000 samples
 Epoch 1/12
 60000/60000 [=====] - 232s 4ms/step - loss: 0.4436 - acc: 0.8613 - val_loss: 0.0557 - val_acc: 0.9809
 Epoch 2/12
 60000/60000 [=====] - 231s 4ms/step - loss: 0.1195 - acc: 0.9660 - val_loss: 0.0399 - val_acc: 0.9879
 Epoch 3/12
 60000/60000 [=====] - 231s 4ms/step - loss: 0.0829 - acc: 0.9772 - val_loss: 0.0301 - val_acc: 0.9902
 Epoch 4/12
 60000/60000 [=====] - 231s 4ms/step - loss: 0.0645 - acc: 0.9817 - val_loss: 0.0299 - val_acc: 0.9911
 Epoch 5/12
 60000/60000 [=====] - 231s 4ms/step - loss: 0.0556 - acc: 0.9842 - val_loss: 0.0258 - val_acc: 0.9921
 Epoch 6/12
 60000/60000 [=====] - 232s 4ms/step - loss: 0.0490 - acc: 0.9864 - val_loss: 0.0233 - val_acc: 0.9928
 Epoch 7/12
 60000/60000 [=====] - 232s 4ms/step - loss: 0.0468 - acc: 0.9872 - val_loss: 0.0229 - val_acc: 0.9929
 Epoch 8/12
 60000/60000 [=====] - 231s 4ms/step - loss: 0.0403 - acc: 0.9887 - val_loss: 0.0185 - val_acc: 0.9950
 Epoch 9/12
 60000/60000 [=====] - 232s 4ms/step - loss: 0.0383 - acc: 0.9896 - val_loss: 0.0192 - val_acc: 0.9941
 Epoch 10/12
 60000/60000 [=====] - 232s 4ms/step - loss: 0.

```
0349 - acc: 0.9902 - val_loss: 0.0256 - val_acc: 0.9919
Epoch 11/12
60000/60000 [=====] - 232s 4ms/step - loss: 0.
0334 - acc: 0.9904 - val_loss: 0.0205 - val_acc: 0.9941
Epoch 12/12
60000/60000 [=====] - 232s 4ms/step - loss: 0.
0313 - acc: 0.9912 - val_loss: 0.0215 - val_acc: 0.9932
```

```
In [7]: # Evaluating the model
score = model5.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

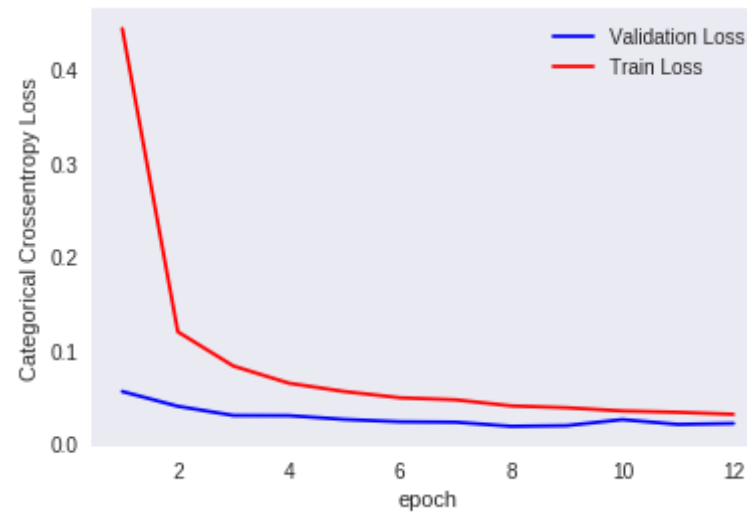
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# Plotting Train and Test Loss VS no. of epochs
# list of epoch numbers
x = list(range(1,epochs+1))

# Validation loss
vy = history.history['val_loss']
# Training loss
ty = history.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)

Test score: 0.021511626279931806
Test accuracy: 0.9932
```



CNN with 7 Convolutional layers and kernel size - (2X2)

```
In [8]: model7 = Sequential()

# Adding first conv layer
model7.add(Conv2D(10, kernel_size=(2, 2),padding='same',activation='relu',input_shape=input_shape))

# Adding second conv layer
model7.add(Conv2D(20, (2, 2), activation='relu'))

# Adding Maxpooling layer
model7.add(MaxPooling2D(pool_size=(3, 3), strides=(1,1)))

# Adding Dropout
model7.add(Dropout(0.3))

# Adding third conv layer
model7.add(Conv2D(40, (2, 2), activation='relu'))
```

```
# Adding Maxpooling layer
model7.add(MaxPooling2D(pool_size=(2, 2),padding='same'))

# Adding fourth conv layer
model7.add(Conv2D(60, (2, 2),padding='same',activation='relu'))

# Adding fifth conv layer
model7.add(Conv2D(120, (2, 2), activation='relu'))

# Adding Maxpooling layer
model7.add(MaxPooling2D(pool_size=(3, 3),padding='same'))

# Adding Dropout
model7.add(Dropout(0.25))

# Adding sixth conv layer
model7.add(Conv2D(120, (2, 2),padding='same',activation='relu'))

# Adding seventh conv layer
model7.add(Conv2D(240, (2, 2), activation='relu'))

# Adding Maxpooling layer
model7.add(MaxPooling2D(pool_size=(2, 2), strides=(1,1)))

# Adding Dropout
model7.add(Dropout(0.25))

# Adding flatten layer
model7.add(Flatten())

# Adding first hidden layer
model7.add(Dense(256, activation='relu',kernel_initializer=he_normal(se
ed=None)))

# Adding Batch Normalization
model7.add(BatchNormalization())
```

```

# Adding Dropout
model7.add(Dropout(0.5))

# Adding second hidden layer
model7.add(Dense(128, activation='relu', kernel_initializer=he_normal(se
ed=None)))

# Adding Dropout
model7.add(Dropout(0.25))

# Adding output layer
model7.add(Dense(num_classes, activation='softmax'))

# Printing model Summary
print(model7.summary())

# Compiling the model
model7.compile(optimizer='adam', loss='categorical_crossentropy', metri
cs=['accuracy'])

# Fitting the data to the model
history = model7.fit(x_train, y_train, batch_size=batch_size, epochs=epoc
hs, verbose=1, validation_data=(x_test, y_test))

```

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 28, 28, 10)	50
conv2d_10 (Conv2D)	(None, 27, 27, 20)	820
max_pooling2d_6 (MaxPooling2	(None, 25, 25, 20)	0
dropout_8 (Dropout)	(None, 25, 25, 20)	0
conv2d_11 (Conv2D)	(None, 24, 24, 40)	3240
max_pooling2d_7 (MaxPooling2	(None, 12, 12, 40)	0
conv2d_12 (Conv2D)	(None, 12, 12, 60)	9660

conv2d_13 (Conv2D)	(None, 11, 11, 120)	28920
max_pooling2d_8 (MaxPooling2D)	(None, 4, 4, 120)	0
dropout_9 (Dropout)	(None, 4, 4, 120)	0
conv2d_14 (Conv2D)	(None, 4, 4, 120)	57720
conv2d_15 (Conv2D)	(None, 3, 3, 240)	115440
max_pooling2d_9 (MaxPooling2D)	(None, 2, 2, 240)	0
dropout_10 (Dropout)	(None, 2, 2, 240)	0
flatten_3 (Flatten)	(None, 960)	0
dense_5 (Dense)	(None, 256)	246016
batch_normalization_2 (Batch Normalization)	(None, 256)	1024
dropout_11 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 128)	32896
dropout_12 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 10)	1290

=====

Total params: 497,076
Trainable params: 496,564
Non-trainable params: 512

None
Train on 60000 samples, validate on 10000 samples
Epoch 1/12
60000/60000 [=====] - 191s 3ms/step - loss: 0.5200 - acc: 0.8264 - val_loss: 0.0519 - val_acc: 0.9848
Epoch 2/12

```

60000/60000 [=====] - 190s 3ms/step - loss: 0.
1031 - acc: 0.9689 - val_loss: 0.0760 - val_acc: 0.9770
Epoch 3/12
60000/60000 [=====] - 189s 3ms/step - loss: 0.
0791 - acc: 0.9770 - val_loss: 0.0325 - val_acc: 0.9903
Epoch 4/12
60000/60000 [=====] - 190s 3ms/step - loss: 0.
0602 - acc: 0.9823 - val_loss: 0.0458 - val_acc: 0.9862
Epoch 5/12
60000/60000 [=====] - 190s 3ms/step - loss: 0.
0535 - acc: 0.9848 - val_loss: 0.0349 - val_acc: 0.9896
Epoch 6/12
60000/60000 [=====] - 190s 3ms/step - loss: 0.
0480 - acc: 0.9860 - val_loss: 0.0196 - val_acc: 0.9935
Epoch 7/12
60000/60000 [=====] - 190s 3ms/step - loss: 0.
0438 - acc: 0.9871 - val_loss: 0.0223 - val_acc: 0.9932
Epoch 8/12
60000/60000 [=====] - 190s 3ms/step - loss: 0.
0369 - acc: 0.9892 - val_loss: 0.0228 - val_acc: 0.9931
Epoch 9/12
60000/60000 [=====] - 189s 3ms/step - loss: 0.
0370 - acc: 0.9895 - val_loss: 0.0296 - val_acc: 0.9917
Epoch 10/12
60000/60000 [=====] - 190s 3ms/step - loss: 0.
0364 - acc: 0.9895 - val_loss: 0.0496 - val_acc: 0.9854
Epoch 11/12
60000/60000 [=====] - 190s 3ms/step - loss: 0.
0324 - acc: 0.9903 - val_loss: 0.0266 - val_acc: 0.9923
Epoch 12/12
60000/60000 [=====] - 189s 3ms/step - loss: 0.
0312 - acc: 0.9908 - val_loss: 0.0352 - val_acc: 0.9909

```

```

In [9]: # Evaluating the model
score = model7.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)

```



```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# Plotting Train and Test Loss VS no. of epochs
# list of epoch numbers
x = list(range(1,epochs+1))

# Validation loss
vy = history.history['val_loss']
# Training loss
ty = history.history['loss']

# Calling the function to draw the plot
plt_dynamic(x, vy, ty)

```

Test score: 0.03518788337475271
Test accuracy: 0.9909



```

In [11]: import pandas as pd
models = pd.DataFrame({'Model': ['CNN 3-conv layer with kernal size (3, 3)', 'CNN 5-conv layer with kernal size(5,5)', 'CNN 7-conv layer with kernal size(2,2)'], 'Test score' : [0.019,0.021,0.035], 'Accuracy': [0.9943,0.9948,0.9909]}, columns = ["Model", "Test score", "Accuracy"])
models

```

Out[11]:

	Model	Test score	Accuracy
0	CNN 3-conv layer with kernal size (3,3)	0.019	0.9943
1	CNN 5-conv layer with kernal size(5,5)	0.021	0.9948
2	CNN 7-conv layer with kernal size(2,2)	0.035	0.9909

Conclusion

- 1.Load mnist dataset.
- 2.Split the dataset into train and test.
- 3.Normalize the data.
- 4.convert class vectors to binary class matrices.
- 5.Implement Softmax classifier with 3 , 5 and 7 CONV layers of of kernel size 3 , 4 and 2 with hidden layers, batchnormalizer and different dropoutrates .
- 6.Ploting Categorical Crossentropy Loss VS No.of Epochs plot .