

Assignment: Different MLP architectures on MNIST dataset

```
In [1]: from keras.utils import np_utils
        from keras.datasets import mnist
        import seaborn as sns
        from keras.initializers import RandomNormal
```

Using TensorFlow backend.

```
In [0]: %matplotlib inline
        import matplotlib.pyplot as plt
        import numpy as np
        import time
        # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
        # https://stackoverflow.com/a/14434334
        # this function is used to update the plots for each epoch and error
        def plt_dynamic(x, vy, ty, ax, colors=['b']):
            ax.plot(x, vy, 'b', label="Validation Loss")
            ax.plot(x, ty, 'r', label="Train Loss")
            plt.legend()
            plt.grid()
            fig.canvas.draw()
```

```
In [3]: # the data, shuffled and split between train and test sets
        (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 1s 0us/step

```
In [4]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
```

```
print("Number of training examples :", X_test.shape[0], "and each image  
is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 3 dimensional vector  
# for each image we have a (28*28) vector  
# we will convert the (28*28) vector into single dimensional vector of  
# 1 * 784  
  
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.sh  
ape[2])  
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2  
])
```

```
In [6]: # after converting the input images from 3d to 2d vectors  
  
print("Number of training examples :", X_train.shape[0], "and each imag  
e is of shape (%d)"%(X_train.shape[1]))  
print("Number of training examples :", X_test.shape[0], "and each image  
is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
In [7]: # An example data point  
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0
```

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	25
4	247	127	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94
0	170	253	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0
2	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253	251	93
3	82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	18	219
0	253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35
0	225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0


```
In [0]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
In [9]: # example data point after normalizing
print(X_train[0])
```

```
[0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.1176471 0.07058824 0.07058824 0.07058824
0.49411765 0.53333333 0.68627451 0.10196078 0.65098039 1.]
```

0.96862745	0.49803922	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588

[illegible]

```
In [10]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ", Y_train[0])
```



```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0.
0. 0.]
```

```
In [0]: from keras.models import Sequential
        from keras.layers import Dense, Activation
```

```
In [0]: # some model parameters

        output_dim = 10
        input_dim = X_train.shape[1]

        batch_size = 128
        nb_epoch = 20
```

MLP + Batch-Norm on 2-hidden Layers + Dropout + AdamOptimizer

```
In [13]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batch-normalization-function-in-keras

        from keras.layers import Dropout
        from keras.layers.normalization import BatchNormalization
        model_drop = Sequential()

        model_drop.add(Dense(450, activation='relu', input_shape=(input_dim,),
        kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
        model_drop.add(BatchNormalization())
        model_drop.add(Dropout(0.5))

        model_drop.add(Dense(108, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
        model_drop.add(BatchNormalization())
        model_drop.add(Dropout(0.5))

        model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_1 (Dense)	(None, 450)	353250
batch_normalization_1 (Batch Normalization)	(None, 450)	1800
dropout_1 (Dropout)	(None, 450)	0
dense_2 (Dense)	(None, 108)	48708
batch_normalization_2 (Batch Normalization)	(None, 108)	432
dropout_2 (Dropout)	(None, 108)	0
dense_3 (Dense)	(None, 10)	1090
=====	=====	=====
Total params: 405,280		
Trainable params: 404,164		
Non-trainable params: 1,116		

```
In [14]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 8s 140us/step - loss: 0.5117 - acc: 0.8455 - val_loss: 0.1711 - val_acc: 0.9443

Epoch 2/20
60000/60000 [=====] - 7s 124us/step - loss: 0.2657 - acc: 0.9211 - val_loss: 0.1310 - val_acc: 0.9590

Epoch 3/20
60000/60000 [=====] - 7s 122us/step - loss: 0.2107 - acc: 0.9360 - val_loss: 0.1172 - val_acc: 0.9647

Epoch 4/20
60000/60000 [=====] - 7s 119us/step - loss: 0.1868 - acc: 0.9443 - val_loss: 0.1019 - val_acc: 0.9684

Epoch 5/20
60000/60000 [=====] - 7s 119us/step - loss: 0.1647 - acc: 0.9505 - val_loss: 0.0913 - val_acc: 0.9723

Epoch 6/20
60000/60000 [=====] - 7s 122us/step - loss: 0.1509 - acc: 0.9540 - val_loss: 0.0869 - val_acc: 0.9738

Epoch 7/20
60000/60000 [=====] - 7s 117us/step - loss: 0.1395 - acc: 0.9569 - val_loss: 0.0831 - val_acc: 0.9748

Epoch 8/20
60000/60000 [=====] - 7s 120us/step - loss: 0.1311 - acc: 0.9598 - val_loss: 0.0859 - val_acc: 0.9734

Epoch 9/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1236 - acc: 0.9624 - val_loss: 0.0778 - val_acc: 0.9759

Epoch 10/20
60000/60000 [=====] - 7s 121us/step - loss: 0.1161 - acc: 0.9651 - val_loss: 0.0701 - val_acc: 0.9781

```

60000/60000 [=====] - 7s 121us/step - loss: 0.
1099 - acc: 0.9664 - val_loss: 0.0727 - val_acc: 0.9773
Epoch 11/20
60000/60000 [=====] - 7s 122us/step - loss: 0.
1053 - acc: 0.9680 - val_loss: 0.0775 - val_acc: 0.9764
Epoch 12/20
60000/60000 [=====] - 7s 119us/step - loss: 0.
1032 - acc: 0.9679 - val_loss: 0.0747 - val_acc: 0.9774
Epoch 13/20
60000/60000 [=====] - 7s 118us/step - loss: 0.
0960 - acc: 0.9700 - val_loss: 0.0706 - val_acc: 0.9787
Epoch 14/20
60000/60000 [=====] - 7s 118us/step - loss: 0.
0932 - acc: 0.9709 - val_loss: 0.0685 - val_acc: 0.9789
Epoch 15/20
60000/60000 [=====] - 7s 120us/step - loss: 0.
0878 - acc: 0.9723 - val_loss: 0.0665 - val_acc: 0.9797
Epoch 16/20
60000/60000 [=====] - 7s 121us/step - loss: 0.
0864 - acc: 0.9729 - val_loss: 0.0636 - val_acc: 0.9810
Epoch 17/20
60000/60000 [=====] - 7s 120us/step - loss: 0.
0794 - acc: 0.9747 - val_loss: 0.0632 - val_acc: 0.9804
Epoch 18/20
60000/60000 [=====] - 7s 121us/step - loss: 0.
0782 - acc: 0.9758 - val_loss: 0.0637 - val_acc: 0.9807
Epoch 19/20
60000/60000 [=====] - 7s 122us/step - loss: 0.
0787 - acc: 0.9752 - val_loss: 0.0596 - val_acc: 0.9808
Epoch 20/20
60000/60000 [=====] - 7s 122us/step - loss: 0.
0727 - acc: 0.9771 - val_loss: 0.0635 - val_acc: 0.9812

```

```

In [15]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

```
# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06352490943533484

Test accuracy: 0.9812



MLP of 2-hidden Layers + AdamOptimizer

```
In [16]: from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization
model_nodrop = Sequential()

model_nodrop.add(Dense(450, activation='relu', input_shape=(input_dim,
), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None
)))

model_nodrop.add(Dense(108, activation='relu', kernel_initializer=Rando
mNormal(mean=0.0, stddev=0.55, seed=None)) )

model_nodrop.add(Dense(output_dim, activation='softmax'))

model_nodrop.summary()
```

Layer (type)	Output Shape	Param #
=====		

dense_4 (Dense)	(None, 450)	353250
dense_5 (Dense)	(None, 108)	48708
dense_6 (Dense)	(None, 10)	1090
=====		
Total params: 403,048		
Trainable params: 403,048		
Non-trainable params: 0		
=====		

```
In [17]: model_nodrop.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
```

```
history = model_nodrop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 100us/step - loss: 0.2367 - acc: 0.9307 - val_loss: 0.1238 - val_acc: 0.9629

Epoch 2/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0915 - acc: 0.9723 - val_loss: 0.1176 - val_acc: 0.9628

Epoch 3/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0603 - acc: 0.9808 - val_loss: 0.0892 - val_acc: 0.9753

Epoch 4/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0435 - acc: 0.9861 - val_loss: 0.1016 - val_acc: 0.9712

Epoch 5/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0332 - acc: 0.9892 - val_loss: 0.0891 - val_acc: 0.9751

Epoch 6/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0261 - acc: 0.9911 - val_loss: 0.0930 - val_acc: 0.9759

Epoch 7/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0264 - acc: 0.9913 - val_loss: 0.0866 - val_acc: 0.9768

```
Epoch 8/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0
174 - acc: 0.9939 - val_loss: 0.0915 - val_acc: 0.9777
Epoch 9/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0
215 - acc: 0.9931 - val_loss: 0.1038 - val_acc: 0.9758
Epoch 10/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0
209 - acc: 0.9925 - val_loss: 0.1200 - val_acc: 0.9731
Epoch 11/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0
184 - acc: 0.9939 - val_loss: 0.1138 - val_acc: 0.9752
Epoch 12/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0
127 - acc: 0.9957 - val_loss: 0.0954 - val_acc: 0.9799
Epoch 13/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0
190 - acc: 0.9938 - val_loss: 0.1217 - val_acc: 0.9744
Epoch 14/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0
115 - acc: 0.9960 - val_loss: 0.1008 - val_acc: 0.9788
Epoch 15/20
60000/60000 [=====] - 6s 100us/step - loss: 0.
0120 - acc: 0.9957 - val_loss: 0.1255 - val_acc: 0.9753
Epoch 16/20
60000/60000 [=====] - 6s 100us/step - loss: 0.
0125 - acc: 0.9958 - val_loss: 0.1000 - val_acc: 0.9793
Epoch 17/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0
115 - acc: 0.9963 - val_loss: 0.1167 - val_acc: 0.9759
Epoch 18/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
144 - acc: 0.9951 - val_loss: 0.1271 - val_acc: 0.9778
Epoch 19/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
106 - acc: 0.9965 - val_loss: 0.1169 - val_acc: 0.9787
Epoch 20/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0
```


064 - acc: 0.9977 - val_loss: 0.1209 - val_acc: 0.9783

```
In [18]: score = model_nodrop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

Test score: 0.12091251760286323
Test accuracy: 0.9783
```



MLP + Batch-Norm on 3-hidden Layers + Dropout + AdamOptimizer

```
In [19]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batch-normalization-function-in-keras

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization
model_drop_3 = Sequential()

model_drop_3.add(Dense(420, activation='relu', input_shape=(input_dim, ), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop_3.add(BatchNormalization())
model_drop_3.add(Dropout(0.5))

model_drop_3.add(Dense(150, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop_3.add(BatchNormalization())
model_drop_3.add(Dropout(0.6))
```

```

model_drop_3.add(Dense(45, activation='relu', kernel_initializer=Random
Normal(mean=0.0, stddev=0.55, seed=None)) )
model_drop_3.add(BatchNormalization())
model_drop_3.add(Dropout(0.7))

model_drop_3.add(Dense(output_dim, activation='softmax'))

model_drop_3.summary()

```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 420)	329700
batch_normalization_3 (Batch Normalization)	(None, 420)	1680
dropout_3 (Dropout)	(None, 420)	0
dense_8 (Dense)	(None, 150)	63150
batch_normalization_4 (Batch Normalization)	(None, 150)	600
dropout_4 (Dropout)	(None, 150)	0
dense_9 (Dense)	(None, 45)	6795
batch_normalization_5 (Batch Normalization)	(None, 45)	180
dropout_5 (Dropout)	(None, 45)	0
dense_10 (Dense)	(None, 10)	460
Total params: 402,565		
Trainable params: 401,335		
Non-trainable params: 1,230		

In [20]: `model_drop_3.compile(optimizer='adam', loss='categorical_crossentropy',`

```
metrics=['accuracy'])
```

```
history = model_drop_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 142us/step - loss: 1.4435 - acc: 0.5397 - val_loss: 0.3629 - val_acc: 0.9110

Epoch 2/20

60000/60000 [=====] - 7s 122us/step - loss: 0.7449 - acc: 0.7655 - val_loss: 0.2329 - val_acc: 0.9336

Epoch 3/20

60000/60000 [=====] - 7s 122us/step - loss: 0.5633 - acc: 0.8330 - val_loss: 0.1909 - val_acc: 0.9441

Epoch 4/20

60000/60000 [=====] - 7s 124us/step - loss: 0.4642 - acc: 0.8677 - val_loss: 0.1632 - val_acc: 0.9530

Epoch 5/20

60000/60000 [=====] - 7s 124us/step - loss: 0.3970 - acc: 0.8895 - val_loss: 0.1486 - val_acc: 0.9569

Epoch 6/20

60000/60000 [=====] - 7s 123us/step - loss: 0.3590 - acc: 0.9024 - val_loss: 0.1374 - val_acc: 0.9606

Epoch 7/20

60000/60000 [=====] - 7s 121us/step - loss: 0.3229 - acc: 0.9143 - val_loss: 0.1278 - val_acc: 0.9652

Epoch 8/20

60000/60000 [=====] - 7s 120us/step - loss: 0.2996 - acc: 0.9209 - val_loss: 0.1230 - val_acc: 0.9656

Epoch 9/20

60000/60000 [=====] - 7s 120us/step - loss: 0.2777 - acc: 0.9262 - val_loss: 0.1210 - val_acc: 0.9687

Epoch 10/20

60000/60000 [=====] - 7s 122us/step - loss: 0.2655 - acc: 0.9330 - val_loss: 0.1125 - val_acc: 0.9686

Epoch 11/20

60000/60000 [=====] - 7s 121us/step - loss: 0.2462 - acc: 0.9367 - val_loss: 0.1103 - val_acc: 0.9712

Epoch 12/20

```

60000/60000 [=====] - 7s 120us/step - loss: 0.
2353 - acc: 0.9394 - val_loss: 0.1048 - val_acc: 0.9741
Epoch 13/20
60000/60000 [=====] - 7s 125us/step - loss: 0.
2263 - acc: 0.9422 - val_loss: 0.1091 - val_acc: 0.9714
Epoch 14/20
60000/60000 [=====] - 8s 130us/step - loss: 0.
2128 - acc: 0.9457 - val_loss: 0.1016 - val_acc: 0.9740
Epoch 15/20
60000/60000 [=====] - 8s 129us/step - loss: 0.
1992 - acc: 0.9483 - val_loss: 0.1040 - val_acc: 0.9736
Epoch 16/20
60000/60000 [=====] - 7s 125us/step - loss: 0.
1973 - acc: 0.9488 - val_loss: 0.0953 - val_acc: 0.9761
Epoch 17/20
60000/60000 [=====] - 7s 125us/step - loss: 0.
1894 - acc: 0.9532 - val_loss: 0.0987 - val_acc: 0.9747
Epoch 18/20
60000/60000 [=====] - 7s 125us/step - loss: 0.
1883 - acc: 0.9531 - val_loss: 0.0918 - val_acc: 0.9770
Epoch 19/20
60000/60000 [=====] - 7s 123us/step - loss: 0.
1800 - acc: 0.9548 - val_loss: 0.0882 - val_acc: 0.9771
Epoch 20/20
60000/60000 [=====] - 7s 121us/step - loss: 0.
1761 - acc: 0.9555 - val_loss: 0.0919 - val_acc: 0.9764

```

```

In [21]: score = model_drop_3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])

```

```
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09186852280043531

Test accuracy: 0.9764



MLP of 3-hidden Layers + AdamOptimizer

```
In [22]: from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization
model_nodrop_3 = Sequential()

model_nodrop_3.add(Dense(420, activation='relu', input_shape=(input_dim,
), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None
)))

model_nodrop_3.add(Dense(150, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )

model_nodrop_3.add(Dense(45, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )

model_nodrop_3.add(Dense(output_dim, activation='softmax'))

model_nodrop_3.summary()
```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 420)	329700
dense_12 (Dense)	(None, 150)	63150
dense_13 (Dense)	(None, 45)	6795
dense_14 (Dense)	(None, 10)	460
Total params: 400,105		
Trainable params: 400,105		
Non-trainable params: 0		

```
In [23]: model_nodrop_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_nodrop_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 102us/step - loss: 0.4762 - acc: 0.8962 - val_loss: 0.2028 - val_acc: 0.9430

Epoch 2/20

60000/60000 [=====] - 5s 91us/step - loss: 0.1408 - acc: 0.9594 - val_loss: 0.1511 - val_acc: 0.9585

Epoch 3/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0950 - acc: 0.9712 - val_loss: 0.1398 - val_acc: 0.9612

Epoch 4/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0762 - acc: 0.9767 - val_loss: 0.1478 - val_acc: 0.9639

Epoch 5/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0601 - acc: 0.9814 - val_loss: 0.1271 - val_acc: 0.9689

Epoch 6/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0485 - acc: 0.9845 - val_loss: 0.1482 - val_acc: 0.9659

Epoch 7/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0463 - acc: 0.9850 - val_loss: 0.1419 - val_acc: 0.9669

Epoch 8/20

60000/60000 [=====] - 6s 96us/step - loss: 0.0352 - acc: 0.9886 - val_loss: 0.1405 - val_acc: 0.9680

Epoch 9/20

60000/60000 [=====] - 6s 100us/step - loss: 0.00349 - acc: 0.9887 - val_loss: 0.1313 - val_acc: 0.9712

Epoch 10/20

60000/60000 [=====] - 6s 94us/step - loss: 0.0362 - acc: 0.9887 - val_loss: 0.1227 - val_acc: 0.9730

Epoch 11/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0275 - acc: 0.9911 - val_loss: 0.1287 - val_acc: 0.9723

Epoch 12/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0220 - acc: 0.9900 - val_loss: 0.1477 - val_acc: 0.9664


```

318 - acc: 0.9900 - val_loss: 0.1477 - val_acc: 0.9684
Epoch 13/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0
257 - acc: 0.9919 - val_loss: 0.1277 - val_acc: 0.9736
Epoch 14/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0
221 - acc: 0.9929 - val_loss: 0.1319 - val_acc: 0.9721
Epoch 15/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0
202 - acc: 0.9934 - val_loss: 0.1432 - val_acc: 0.9743
Epoch 16/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0
309 - acc: 0.9912 - val_loss: 0.1229 - val_acc: 0.9756
Epoch 17/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0
167 - acc: 0.9949 - val_loss: 0.1499 - val_acc: 0.9699
Epoch 18/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0
143 - acc: 0.9954 - val_loss: 0.1324 - val_acc: 0.9733
Epoch 19/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0
208 - acc: 0.9934 - val_loss: 0.1676 - val_acc: 0.9721
Epoch 20/20
60000/60000 [=====] - 5s 92us/step - loss: 0.0
163 - acc: 0.9950 - val_loss: 0.1399 - val_acc: 0.9747

```

```

In [24]: score = model_nodrop_3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epo

```

```

chs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter vali
dation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal
to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.13994209184171222

Test accuracy: 0.9747



MLP + Batch-Norm on 5-hidden Layers + Dropout + AdamOptimizer

```
In [0]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batch-normalization-function-in-keras
```

```
from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization
model_drop_5 = Sequential()

model_drop_5.add(Dense(512, activation='relu', input_shape=(input_dim, ), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop_5.add(BatchNormalization())
model_drop_5.add(Dropout(0.8))

model_drop_5.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop_5.add(BatchNormalization())
model_drop_5.add(Dropout(0.7))

model_drop_5.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop_5.add(BatchNormalization())
model_drop_5.add(Dropout(0.6))

model_drop_5.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop_5.add(BatchNormalization())
model_drop_5.add(Dropout(0.5))

model_drop_5.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop_5.add(BatchNormalization())
model_drop_5.add(Dropout(0.4))

model_drop_5.add(Dense(output_dim, activation='softmax'))
```

```
In [0]: model_drop_5.summary()
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
dense_15 (Dense)                (None, 512)                401920
batch_normalization_6 (Batch Normalization) (None, 512)                2048
dropout_6 (Dropout)              (None, 512)                 0
dense_16 (Dense)                (None, 256)               131328
batch_normalization_7 (Batch Normalization) (None, 256)                1024
dropout_7 (Dropout)              (None, 256)                 0
dense_17 (Dense)                (None, 128)               32896
batch_normalization_8 (Batch Normalization) (None, 128)                 512
dropout_8 (Dropout)              (None, 128)                 0
dense_18 (Dense)                (None, 64)                8256
batch_normalization_9 (Batch Normalization) (None, 64)                 256
dropout_9 (Dropout)              (None, 64)                  0
dense_19 (Dense)                (None, 32)                2080
batch_normalization_10 (Batch Normalization) (None, 32)                 128
dropout_10 (Dropout)             (None, 32)                  0
dense_20 (Dense)                (None, 10)                 330
=====
Total params: 580,778
Trainable params: 578,794
Non-trainable params: 1,984
=====

```

```
In [26]: model_drop_5.compile(optimizer='adam', loss='categorical_crossentropy',
```

```
metrics=['accuracy'])
```

```
history = model_drop_5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 13s 209us/step - loss: 2.4181 - acc: 0.1676 - val_loss: 1.9354 - val_acc: 0.4156

Epoch 2/20

60000/60000 [=====] - 11s 182us/step - loss: 1.8200 - acc: 0.3338 - val_loss: 1.1060 - val_acc: 0.6734

Epoch 3/20

60000/60000 [=====] - 11s 187us/step - loss: 1.3506 - acc: 0.4993 - val_loss: 0.8121 - val_acc: 0.7470

Epoch 4/20

60000/60000 [=====] - 11s 183us/step - loss: 1.1110 - acc: 0.5875 - val_loss: 0.6596 - val_acc: 0.7935

Epoch 5/20

60000/60000 [=====] - 11s 181us/step - loss: 0.9731 - acc: 0.6512 - val_loss: 0.5571 - val_acc: 0.8273

Epoch 6/20

60000/60000 [=====] - 11s 183us/step - loss: 0.8757 - acc: 0.6937 - val_loss: 0.4738 - val_acc: 0.8666

Epoch 7/20

60000/60000 [=====] - 11s 183us/step - loss: 0.7992 - acc: 0.7297 - val_loss: 0.4131 - val_acc: 0.8917

Epoch 8/20

60000/60000 [=====] - 11s 186us/step - loss: 0.7349 - acc: 0.7582 - val_loss: 0.3691 - val_acc: 0.9019

Epoch 9/20

60000/60000 [=====] - 11s 184us/step - loss: 0.6883 - acc: 0.7830 - val_loss: 0.3261 - val_acc: 0.9145

Epoch 10/20

60000/60000 [=====] - 11s 183us/step - loss: 0.6301 - acc: 0.8059 - val_loss: 0.3023 - val_acc: 0.9180

Epoch 11/20

60000/60000 [=====] - 11s 183us/step - loss: 0.5869 - acc: 0.8256 - val_loss: 0.2608 - val_acc: 0.9299

Epoch 12/20

```

60000/60000 [=====] - 11s 184us/step - loss:
0.5497 - acc: 0.8404 - val_loss: 0.2367 - val_acc: 0.9353
Epoch 13/20
60000/60000 [=====] - 11s 183us/step - loss:
0.5190 - acc: 0.8523 - val_loss: 0.2205 - val_acc: 0.9409
Epoch 14/20
60000/60000 [=====] - 11s 186us/step - loss:
0.4923 - acc: 0.8640 - val_loss: 0.2160 - val_acc: 0.9426
Epoch 15/20
60000/60000 [=====] - 11s 187us/step - loss:
0.4735 - acc: 0.8736 - val_loss: 0.2039 - val_acc: 0.9459
Epoch 16/20
60000/60000 [=====] - 11s 184us/step - loss:
0.4423 - acc: 0.8818 - val_loss: 0.1992 - val_acc: 0.9473
Epoch 17/20
60000/60000 [=====] - 11s 183us/step - loss:
0.4267 - acc: 0.8888 - val_loss: 0.1888 - val_acc: 0.9511
Epoch 18/20
60000/60000 [=====] - 11s 183us/step - loss:
0.4148 - acc: 0.8916 - val_loss: 0.1823 - val_acc: 0.9540
Epoch 19/20
60000/60000 [=====] - 11s 182us/step - loss:
0.3932 - acc: 0.8989 - val_loss: 0.1781 - val_acc: 0.9551
Epoch 20/20
60000/60000 [=====] - 12s 196us/step - loss:
0.3867 - acc: 0.9021 - val_loss: 0.1772 - val_acc: 0.9543

```

```

In [27]: score = model_drop_5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])

```

```
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1772317446306348

Test accuracy: 0.9543



MLP of 5-hidden Layers + AdamOptimizer

```
In [28]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batch-normalization-function-in-keras

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization
from keras.initializers import he_normal
model_nodrop_5 = Sequential()

model_nodrop_5.add(Dense(512, activation='relu', input_shape=(input_dim, ), kernel_initializer=he_normal( seed=None)))

model_nodrop_5.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)) )

model_nodrop_5.add(Dense(128, activation='relu', kernel_initializer=he_normal(seed=None)) )

model_nodrop_5.add(Dense(64, activation='relu', kernel_initializer=he_normal(seed=None)) )

model_nodrop_5.add(Dense(32, activation='relu', kernel_initializer=he_normal(seed=None)) )

model_nodrop_5.add(Dense(output_dim, activation='softmax'))
model_nodrop_5.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_21 (Dense)	(None, 512)	401920
dense_22 (Dense)	(None, 256)	131328
dense_23 (Dense)	(None, 128)	32896
dense_24 (Dense)	(None, 64)	8256
dense_25 (Dense)	(None, 32)	2080

dense_26 (Dense)	(None, 10)	330
------------------	------------	-----

```

Total params: 576,810
Trainable params: 576,810
Non-trainable params: 0

```

```

In [29]: model_nodrop_5.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

history = model_nodrop_5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 9s 151us/step - loss: 0.2507 - acc: 0.9238 - val_loss: 0.1110 - val_acc: 0.9677
Epoch 2/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0908 - acc: 0.9722 - val_loss: 0.1099 - val_acc: 0.9668
Epoch 3/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0612 - acc: 0.9810 - val_loss: 0.1000 - val_acc: 0.9708
Epoch 4/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0460 - acc: 0.9854 - val_loss: 0.0849 - val_acc: 0.9779
Epoch 5/20
60000/60000 [=====] - 8s 137us/step - loss: 0.0379 - acc: 0.9879 - val_loss: 0.0842 - val_acc: 0.9773
Epoch 6/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0288 - acc: 0.9907 - val_loss: 0.0949 - val_acc: 0.9745
Epoch 7/20
60000/60000 [=====] - 8s 138us/step - loss: 0.0297 - acc: 0.9908 - val_loss: 0.0709 - val_acc: 0.9819
Epoch 8/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0190 - acc: 0.9938 - val_loss: 0.0953 - val_acc: 0.9778
Epoch 9/20

```

```

60000/60000 [=====] - 8s 138us/step - loss: 0.
0229 - acc: 0.9927 - val_loss: 0.0770 - val_acc: 0.9791
Epoch 10/20
60000/60000 [=====] - 8s 136us/step - loss: 0.
0188 - acc: 0.9937 - val_loss: 0.0922 - val_acc: 0.9793
Epoch 11/20
60000/60000 [=====] - 8s 137us/step - loss: 0.
0179 - acc: 0.9943 - val_loss: 0.0925 - val_acc: 0.9781
Epoch 12/20
60000/60000 [=====] - 8s 136us/step - loss: 0.
0171 - acc: 0.9944 - val_loss: 0.0781 - val_acc: 0.9807
Epoch 13/20
60000/60000 [=====] - 8s 137us/step - loss: 0.
0135 - acc: 0.9957 - val_loss: 0.0831 - val_acc: 0.9787
Epoch 14/20
60000/60000 [=====] - 8s 136us/step - loss: 0.
0130 - acc: 0.9960 - val_loss: 0.0957 - val_acc: 0.9805
Epoch 15/20
60000/60000 [=====] - 8s 133us/step - loss: 0.
0133 - acc: 0.9958 - val_loss: 0.0770 - val_acc: 0.9809
Epoch 16/20
60000/60000 [=====] - 9s 143us/step - loss: 0.
0133 - acc: 0.9959 - val_loss: 0.0942 - val_acc: 0.9765
Epoch 17/20
60000/60000 [=====] - 9s 143us/step - loss: 0.
0136 - acc: 0.9956 - val_loss: 0.0870 - val_acc: 0.9821
Epoch 18/20
60000/60000 [=====] - 8s 133us/step - loss: 0.
0122 - acc: 0.9963 - val_loss: 0.0998 - val_acc: 0.9787
Epoch 19/20
60000/60000 [=====] - 8s 136us/step - loss: 0.
0082 - acc: 0.9974 - val_loss: 0.1091 - val_acc: 0.9765
Epoch 20/20
60000/60000 [=====] - 9s 143us/step - loss: 0.
0118 - acc: 0.9962 - val_loss: 0.0982 - val_acc: 0.9805

```

```

In [30]: score = model_nodrop_5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

```

```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09818555498997499

Test accuracy: 0.9805



MLP with Batch Normalization and Dropout

```
In [33]: import pandas as pd
models = pd.DataFrame({'Hidden layers': ['2', '3', "5"], 'Test score' :
[0.06,0.09,0.17], 'Accuracy': [0.9812,0.9764,0.9543]}, columns = ["Hidd
en layers", "Test score", "Accuracy"])
models
```

Out[33]:

	Hidden layers	Test score	Accuracy
0	2	0.06	0.9812
1	3	0.09	0.9764
2	5	0.17	0.9543

MLP without Batch Normalization and Dropout

```
In [34]: models = pd.DataFrame({'Hidden layers': ['2', '3', "5"], 'Test score' :
[0.12,0.13,0.09], 'Accuracy': [0.9783,0.9747,0.9805]}, columns = ["Hidd
```

```
en_layers","Test score","Accuracy"])\nmodels
```

Out[34]:

	Hidden layers	Test score	Accuracy
0	2	0.12	0.9783
1	3	0.13	0.9747
2	5	0.09	0.9805

Conclusion

- 1.Load mnist dataset.
- 2.Split the dataset into train and test.
- 3.converting the input images from 3d to 2d vectors.
- 4.Normalize the data.
- 5.Implement Softmax classifier with 2 , 3 and 5 hidden layers with batchnormalizer and different dropoutrates.
- 6.Implement Softmax classifier with 2 , 3 and 5 hidden layers without batchnormalizer and dropoutrates.
- 7.Ploting Categorical Crossentropy Loss VS No.of Epochs plot .